

Integrated Project: Maji Ndogo Part 2

- **Clustering data to unveil Maji Ndogo's water crisis**

Cleaning our data

Ok, bring up the employee table. It has info on all of our workers, but note that the email addresses have not been added. We will have to send them reports and figures, so let's update it. Luckily the emails for our department are easy:

first_name.last_name@ndogowater.gov.

We can determine the email address for each employee by:

- selecting the employee_name column
- replacing the space with a full stop
- make it lowercase
- and stitch it all together

- Replace the space with a full stop AND Make it all lowercase PAGE (6)

SELECT

LOWER(REPLACE(employee_name, ' ', '.'))

FROM employee;

Result Grid		Filter Rows:
	LOWER(REPLACE(employee_name,' ','.'))	
▶	amara.jengo	
	bello.azibo	
	bakari.iniko	
	malachi.mavuso	
	cheche.buhle	
	zuriel.matembo	
	deka.osumare	
	lalitha.kaburi	
	enitan.zuri	
	farai.nia	
	gamba.shani	
	harith.nyota	
	isoke.amani	
	jengo.tumaini	
	luntu.acha	

- Replace the space with a full stop AND Make it all lower case,
- and use CONCAT() '@ndogowater.gov') AS new_email -- add it all together PAGE (6)

SELECT

CONCAT(

LOWER(REPLACE(employee_name,' ','.')), '@ndogowater.gov') AS new_email

FROM employee;

Result Grid	Filter Rows:	Export:	Wrap
new_email			
amara.jengo@ndogowater.gov			
belo.azibo@ndogowater.gov			
bakari.iniko@ndogowater.gov			
malachi.mavuso@ndogowater.gov			
cheche.buhle@ndogowater.gov			
Result 7			
Output			

- Use the employee table to count how many of our employees live in each town.
- Think carefully about what function we should use and how we should aggregate the data.

SELECT DISTINCT

province_name,

town_name,

COUNT(town_name) AS num_employees

FROM employee

GROUP BY

1,2

ORDER BY

1,3 DESC;

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
town_name	province_name	num_employee	
Dahabu	Amanzi	6	
Harare	Akatsi	3	
Harare	Kilimani	2	
Ilanga	Sokoto	2	
Ilanga	Kilimani	1	
Result 10			
Output			

- Let's first look at the number of records each employee collected. So find the correct table,
- figure out what function to use and how to group, order
- and limit the results to only see the top 3 employee_ids
- with the highest number of locations visited PAGE (10)

SELECT

assigned_employee_id,

COUNT(visit_count) AS number_of_visits

FROM


visits


GROUP BY

assigned_employee_id

ORDER BY

number_of_visits DESC

Result Grid  Filter Rows: <input type="text"/>		
	assigned_employee_id	number_of_visits
▶	1	3708
	30	3676
	34	3539
	3	3420
	10	3407
	8	3351
	5	3284
	36	3249
	48	2933
	28	2762
	12	2561
	42	2496
	40	2344
	38	2121
	2	2033
	74	2015

Result 2 x 

- Create a query that counts the number of records per town PAGE (11)

SELECT DISTINCT

town_name,

COUNT(town_name) OVER (PARTITION **BY** town_name) **AS** records_per_town

FROM

Location;

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
town_name	records_per_town		
Abidjan	400		
Amara	710		
Amina	1090		
Asmara	930		
Bahari	470		

Result 5 x

Output

- Now count the records per province. PAGE (12)

SELECT DISTINCT

province_name,

COUNT(province_name) OVER (PARTITION **BY** province_name) **AS** records_per_province

FROM

location;

Result Grid	Filter Rows:	Export:	Wrap Cell Content:
province_name	records_per_province		
Kilimani	9510		
Akatsi	8940		
Sokoto	8220		
Amanzi	6950		
Hawassa	6030		

Result 6 x

Output

- Create a result set showing:
 - province_name
 - town_name
 - An aggregated count of records for each town (consider naming this records_per_town).
 - Ensure your data is grouped by both province_name and town_name.
- 2. Order your results primarily by province_name. Within each province, further sort the towns by their record counts in descending order.

SELECT DISTINCT

province_name,

town_name,

COUNT(town_name) AS records_per_town

FROM

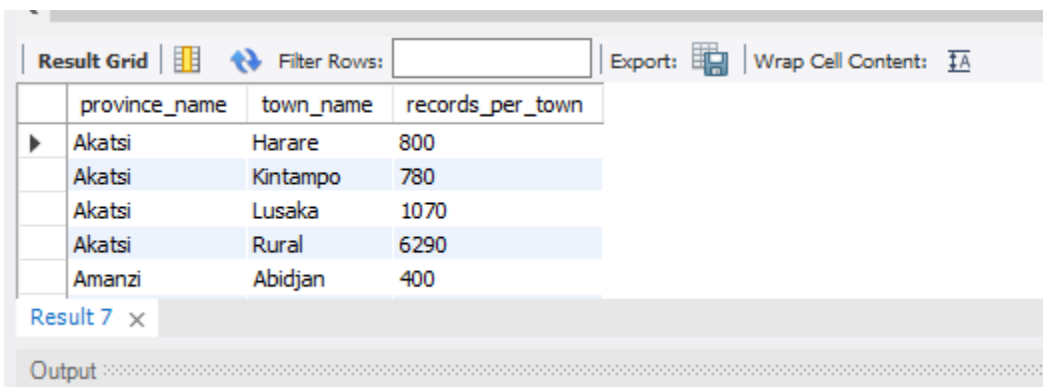
location

GROUP BY

province_name,

town_name

ORDER BY province_name;



The screenshot shows a database query result grid with the following data:

	province_name	town_name	records_per_town
▶	Akatsi	Harare	800
	Akatsi	Kintampo	780
	Akatsi	Lusaka	1070
	Akatsi	Rural	6290
	Amanzi	Abidjan	400

Below the table, there is a tab labeled "Result 7" and an "Output" section.

- Finally, look at the number of records for each location type PAGE (13)

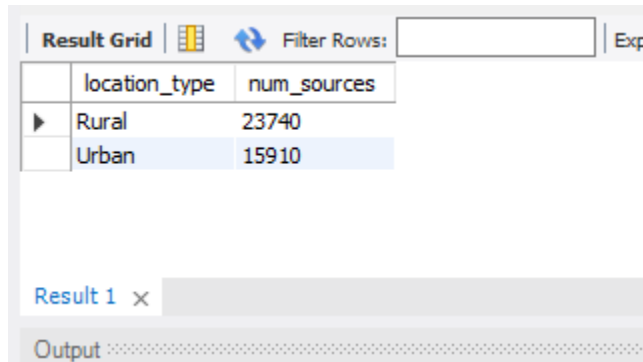
SELECT DISTINCT

location_type,

COUNT(location_type) OVER (PARTITION BY location_type) AS num_sources

FROM

location;



The screenshot shows a 'Result Grid' interface with a 'Filter Rows' search bar. The grid contains two columns: 'location_type' and 'num_sources'. There are two rows: 'Rural' with a value of 23740, and 'Urban' with a value of 15910. Below the grid, there is a tab labeled 'Result 1' and an 'Output' section.

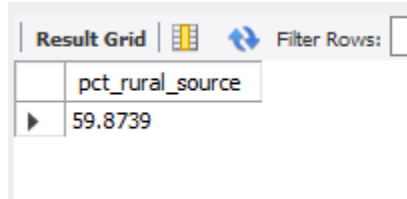
	location_type	num_sources
▶	Rural	23740
	Urban	15910

- We can see that there are more rural sources than urban,
- but it's really hard to understand those numbers. Percentages are more relatable.
- If we use SQL as a very overpowered calculator:

SELECT

$23740 / (15910 + 23740) * 100$ AS pct_rural_source

We can see that 60% of all water sources in the data set are in rural communities.



The screenshot shows a 'Result Grid' interface with a 'Filter Rows' search bar. The grid contains one column: 'pct_rural_source'. There is one row with a value of 59.8739.

	pct_rural_source
▶	59.8739

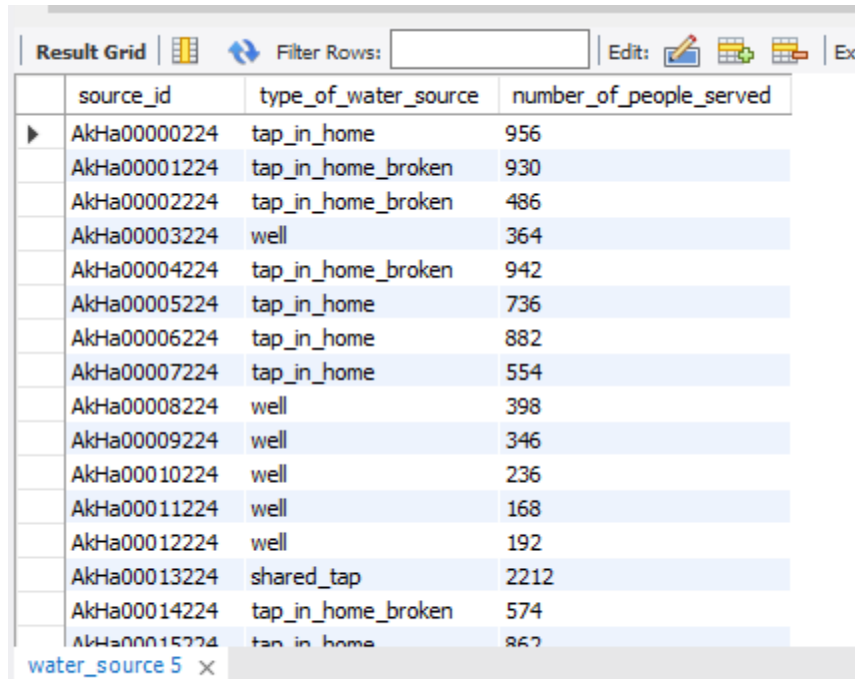
So again, what are some of the insights we gained from the location table?

1. Our entire country was properly canvassed, and our dataset represents the situation on the ground.
2. 60% of our water sources are in rural communities across Maji Ndogo. We need to keep this in mind when we make decisions

SELECT * FROM

water_source

LIMIT 50;



The screenshot shows a database query result grid with the following columns: source_id, type_of_water_source, and number_of_people_served. The grid contains 17 rows of data. The interface includes a 'Result Grid' tab, a 'Filter Rows' search bar, and an 'Edit' button. A tab labeled 'water_source 5' is visible at the bottom left.

	source_id	type_of_water_source	number_of_people_served
▶	AkHa00000224	tap_in_home	956
	AkHa00001224	tap_in_home_broken	930
	AkHa00002224	tap_in_home_broken	486
	AkHa00003224	well	364
	AkHa00004224	tap_in_home_broken	942
	AkHa00005224	tap_in_home	736
	AkHa00006224	tap_in_home	882
	AkHa00007224	tap_in_home	554
	AkHa00008224	well	398
	AkHa00009224	well	346
	AkHa00010224	well	236
	AkHa00011224	well	168
	AkHa00012224	well	192
	AkHa00013224	shared_tap	2212
	AkHa00014224	tap_in_home_broken	574
	AkHa00015224	tap_in_home	862

We have access to different water source types and the number of people using each source. These are the questions that I am curious about.

1. How many people did we survey in total?
2. How many wells, taps and rivers are there?
3. How many people share particular types of water sources on average?
4. How many people are getting water from each type of source?

1. How many people did we survey in total?

SELECT

SUM(number_of_people_served) AS SURVEY_IN_TOTAL -- > EQUAL = '27628140'

FROM

water_source;

Result Grid	Filter Rows:
total_of_people_served	
27628140	

Result 20 x

2. How many wells, taps and rivers are there? PAGE(15)

SELECT DISTINCT

type_of_water_source,

**COUNT(type_of_water_source) OVER (PARTITION BY type_of_water_source) AS
number_of_sources**

FROM

water_source;

Result Grid	Filter Rows:
type_of_water_source	number_of_sources
river	3379
shared_tap	5767
tap_in_home	7265
tap_in_home_broken	5856
well	17383

3. How many people share particular types of water sources on average? PAGE(16)



SELECT DISTINCT

type_of_water_source,

**ROUND(AVG(number_of_people_served) OVER (PARTITION BY
type_of_water_source)) AS number_of_sources**

FROM

water_source;

Result Grid   Filter Rows: <input type="text"/> Export:		
	type_of_water_source	number_of_sources
▶	river	699
	shared_tap	2071
	tap_in_home	644
	tap_in_home_broken	649
	well	279

- Now let's calculate the total number of people served by each type of water source in total, to make it easier to interpret, order them so the most people served by a source is at the top. PAGE(19)



SELECT DISTINCT


type_of_water_source,

**ROUND(SUM(number_of_people_served) OVER (PARTITION BY
type_of_water_source)) AS population_served**

FROM

water_source;

Result Grid   Filter Rows: <input type="text"/> Export:		
	type_of_water_source	population_served
▶	river	2362544
	shared_tap	11945272
	tap_in_home	4678880
	tap_in_home_broken	3799720
	well	4841724

Result 37 x 

Next, calculate the percentages using the total we just got.

Let's round that off to 0 decimals, and order the results. PAGE(21)

SELECT DISTINCT

type_of_water_source,

**FORMAT((SUM(number_of_people_served)/27628140)*100,0) AS
percentage_people_per_source**

FROM

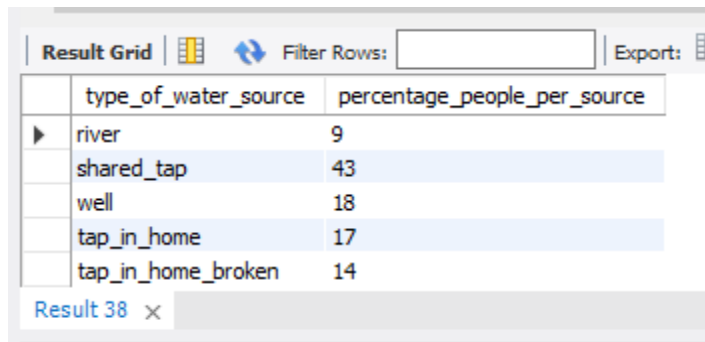
water_source

GROUP BY

type_of_water_source

ORDER BY

percentage_people_per_source DESC;



The screenshot shows a 'Result Grid' window with a table of water sources. The table has two columns: 'type_of_water_source' and 'percentage_people_per_source'. The rows are ordered from highest to lowest percentage. The 'tap_in_home' row is highlighted in blue. Below the table, there is a tab labeled 'Result 38' with a close button 'x'.

	type_of_water_source	percentage_people_per_source
▶	river	9
	shared_tap	43
	well	18
	tap_in_home	17
	tap_in_home_broken	14

Result 38 x

**So use a window function on the total people served column, converting it into a rank.
PAGE(23)**

- But think about this: If someone has a tap in their home,
- they already have the best source available. Since we can't do anything more to improve
- this, we should remove tap_in_home from the ranking before we continue.

SELECT DISTINCT

type_of_water_source,

SUM(number_of_people_served) AS people_served,

**RANK() OVER (ORDER BY SUM(number_of_people_served) DESC) AS
rank_by_population**

FROM

water_source





WHERE

type_of_water_source != 'tap_in_home'

GROUP BY

type_of_water_source

LIMIT 50;

Result Grid   Filter Rows: <input type="text"/> Export:  			
	type_of_water_source	people_served	rank_by_population
▶	shared_tap	11945272	1
	well	4841724	2
	tap_in_home_broken	3799720	3
	river	2362544	4

So create a query to do this, and keep these requirements in mind: PAGE(24)

1. The sources within each type should be assigned a rank.
2. Limit the results to only improvable sources.
3. Think about how to partition, filter and order the results set.
4. Order the results to see the top of the list.

SELECT DISTINCT

source_id,

type_of_water_source,

number_of_people_served,

**DENSE_RANK() OVER (ORDER BY number_of_people_served DESC) AS
priority_rank**

FROM

water_source

WHERE

type_of_water_source != 'tap_in_home'

LIMIT 50;

Result Grid Filter Rows: Export: Wrap Cell Content:				
	source_id	type_of_water_source	number_of_people_served	priority_rank
▶	HaRu19509224	shared_tap	3998	1
	AkRu05603224	shared_tap	3998	1
	AkRu04862224	shared_tap	3996	2
	KiHa22867224	shared_tap	3996	2
	AmAs10911224	shared_tap	3996	2
	HaRu19839224	shared_tap	3994	3
	KiZu31330224	shared_tap	3994	3
	KiRu28630224	shared_tap	3992	4
	KiZu31415224	shared_tap	3992	4
	KiRu26218224	shared_tap	3990	5
	AmPw12653224	shared_tap	3990	5
	SoRu38511224	shared_tap	3990	5
	SoRu35346224	shared_tap	3988	6
	SoRu36081224	shared_tap	3988	6
	SoRu36687224	shared_tap	3988	6
	HaRu18763224	shared_tap	3988	6

Result 9 x

Analysing queues

Ok, these are some of the things I think are worth looking at:

1. How long did the survey take?
2. What is the average total queue time for water?
3. What is the average queue time on different days?
4. How can we communicate this information efficiently?

1. Question 1: PAGE(27)
 - To calculate how long the survey took,
 - we need to get the first and last dates (which functions can find the largest/smallest value), and subtract
 - them. Remember with DateTime data,
 - we can't just subtract the values. We have to use a function to get the difference in days.

SELECT

TIMESTAMPDIFF(DAY,MIN(time_of_record),MAX(time_of_record)) AS DURATION

FROM

visits;

Result Grid	Filter Rows:
DURATION	
▶ 924	

2. Question 2: PAGE(28)

- Let's see how long people have to queue on average in Maji Ndogo.
- Keep in mind that many sources like taps_in_home have no queues. These
- are just recorded as 0 in the time_in_queue column,
- so when we calculate averages, we need to exclude those rows. Try using NULLIF() to do this.

SELECT

AVG(NULLIF(time_in_queue,0)) AS AVG_time_in_queue

FROM

visits;

- ☐ You should get a queue time of about 123 min. So on average,
- ☐ People take two hours to fetch water if they don't have a tap in their homes.

Result Grid	Filter Rows:
AVG_time_in_queue	
▶ 123.2574	

**So let's look at the queue times aggregated across the different days of the week.
PAGE(29)**

SELECT DISTINCT

DAYNAME(time_of_record) AS day_of_week,

ROUND(AVG(time_in_queue)) AS avg_queue_time

FROM

visits

GROUP BY day_of_week

ORDER BY day_of_week;

Result Grid			Filter Rows:
	day_of_week	avg_queue_time	
►	Friday	53	
	Monday	60	
	Saturday	246	
	Sunday	82	
	Thursday	46	
	Tuesday	47	
	Wednesday	43	

4. Question 4: PAGE(31)

- We can also look at what time during the day people collect water.
- Try to order the results in a meaningful way.

SELECT

TIME_FORMAT(TIME(time_of_record),'%H:00') AS hour_of_day,

ROUND(AVG(time_in_queue)) AS avg_queue_time

FROM

visits

GROUP BY

hour_of_day

ORDER BY

avg_queue_time;

- Can you see that mornings and evenings are the busiest?
- It looks like people collect water before and after work.
- Wouldn't it be nice to break down the queue times for each hour of each day? In a spreadsheet,
- we can just create a pivot table.

Result Grid		
		Filter Rows:
	hour_of_day	avg_queue_time
▶	11:00	46
	12:00	47
	13:00	47
	14:00	47
	16:00	47
	10:00	48
	15:00	48
	09:00	49
	18:00	147
	07:00	149
	08:00	149
	17:00	149
	06:00	149
	19:00	168

SELECT

TIME_FORMAT(TIME(time_of_record), '%H:00') AS hour_of_day,

DAYNAME(time_of_record) AS DAY_NAME,

CASE

WHEN DAYNAME(time_of_record) = 'Sunday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Monday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Saturday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Tuesday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Wednesday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Thursday'

THEN time_in_queue

WHEN DAYNAME(time_of_record) = 'Friday'


```

        THEN time_in_queue

ELSE NULL

END AS Time_of_waiting

FROM

    mdd_water_services.visits

WHERE

    time_in_queue != 0

    • This excludes other sources with 0 queue times.

LIMIT 50;

```

Result Grid			
Filter Rows:			
	hour_of_day	DAY_NAME	Time_of_waiting
▶	09:00	Friday	15
	09:00	Friday	62
	10:00	Friday	28
	10:00	Friday	9
	11:00	Friday	17
	11:00	Friday	240
	12:00	Friday	20
	12:00	Friday	171
	13:00	Friday	28
	13:00	Friday	16
	13:00	Friday	56
	14:00	Friday	11
	14:00	Friday	24
	15:00	Friday	211
	16:00	Friday	16
	07:00	Saturday	107

- Creating a pivot table of time waiting in each day PAGE(35)

```

SELECT

TIME_FORMAT(TIME(time_of_record), '%H:00') AS hour_of_day,

-- Sunday

ROUND(AVG(

    CASE

```

```

        WHEN DAYNAME(time_of_record) = 'Sunday'

        THEN time_in_queue

    ELSE NULL

END),o) AS Sunday,
-- Monday
ROUND(AVG(

    CASE

        WHEN DAYNAME(time_of_record) = 'Monday'

        THEN time_in_queue

    ELSE NULL

        END),o) AS Monday,
-- Tuesday
ROUND(AVG(

    CASE

        WHEN DAYNAME(time_of_record) = 'Tuesday'

        THEN time_in_queue

    ELSE NULL

        END),o) AS Tuesday,
-- Wednesday
ROUND(AVG(

    CASE

        WHEN DAYNAME(time_of_record) = 'Wednesday'

        THEN time_in_queue

    ELSE NULL

        END),o) AS Wednesday,
-- Thursday
ROUND(AVG(

```

```

CASE
WHEN DAYNAME(time_of_record) = 'Thursday'
THEN time_in_queue
ELSE NULL
END),o) AS Thursday,
-- Friday
ROUND(AVG(
CASE
WHEN DAYNAME(time_of_record) = 'Friday'
THEN time_in_queue
ELSE NULL
END),o) AS Friday,
-- Saturday
ROUND(AVG(
CASE
WHEN DAYNAME(time_of_record) = 'Saturday'
THEN time_in_queue
ELSE NULL
END),o) AS Saturday
FROM
visits
WHERE
time_in_queue != 0 -- this excludes other sources with 0 queue times
GROUP BY
hour_of_day
ORDER BY
hour_of_day;

```

<div> <div>Result Grid</div> <div> </div> <div> Filter Rows: <input type="text"/> </div> <div> Export: </div> <div> Wrap Cell Content: </div> </div>								
	hour_of_day	Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
▶	06:00	79	190	134	112	134	153	247
	07:00	82	186	128	111	139	156	247
	08:00	86	183	130	119	129	153	247
	09:00	84	127	105	94	99	107	252
	10:00	83	119	99	89	95	112	259
	11:00	78	115	102	86	99	104	236
	12:00	78	115	97	88	96	109	239
	13:00	81	122	97	98	101	115	242
	14:00	83	127	104	92	96	110	244
	15:00	83	126	104	88	92	110	248
	16:00	83	127	99	90	99	109	251
	17:00	79	181	135	121	129	151	251
	18:00	80	174	122	113	132	158	240
	19:00	127	159	145	176	137	103	282